
COMPUTER ARCHITECTURE

LAB THREE

PREPARED BY

AHMED ALY GAMAL EL-DIN EL-GHANNAM

ELECTRONICS AND COMMUNICATIONS - LEVEL 4
ID: 19015292

YAHIA WALID EL-DAKHAKHNY

ELECTRONICS AND COMMUNICATIONS - LEVEL 4
ID: 19016891

MAY 2024

Contents

1	Introduction	1
2	Prerequisites	2
2.1	Host System Specifications	2
2.2	Installation Process	2
2.3	Compilation Settings	3
3	CPU Version of The Program	4
3.1	Program Code	4
3.2	Runtime & Profiling Results	7
3.3	Observation	9
4	GPU Version of The Program	10
4.1	Program Code	10
4.2	Runtime & Profiling Results	14
4.3	Observation	18

List of Figures

1	Host System Specifications	2
2	Runtime at N = 50	7
3	Runtime at N = 100	7
4	Runtime at N = 500	7
5	Function Timing at N = 500 (gprof)	7
6	Runtime at N = 1000	7
7	Function Timing at N = 1000 (gprof)	8
8	Runtime at N = 2000	8
9	Function Timing at N = 2000 (gprof)	8
10	Runtime at N = 5000	8
11	Function Timing at N = 5000 (gprof)	8
12	CPU – Runtime vs N	9
13	Program Timing Information at N = 50 (Thread Count of 16)	14
14	Program Timing Information at N = 100 (Thread Count of 16)	14
15	Program Timing Information at N = 500 (Thread Count of 16)	14
16	Program Timing Information at N = 1000 (Thread Count of 16)	15
17	Program Timing Information at N = 2000 (Thread Count of 16)	15
18	Program Timing Information at N = 5000 (Thread Count of 16)	15
19	Program Timing Information at N = 50 (Thread Count of 32)	16
20	Program Timing Information at N = 100 (Thread Count of 32)	16
21	Program Timing Information at N = 500 (Thread Count of 32)	16
22	Program Timing Information at N = 1000 (Thread Count of 32)	17
23	Program Timing Information at N = 2000 (Thread Count of 32)	17
24	Program Timing Information at N = 5000 (Thread Count of 32)	17
25	GPU – Runtime vs N (Thread Count of 16)	18
26	GPU – Runtime vs N (Thread Count of 32)	18

Listings

1	Included Libraries & Macro	4
2	Function to Populate Matrices with Random Values	4
3	Function to Print Matrix (For Testing Only)	4
4	Function to Add Two Matrices	5
5	unction to Multiply Two Matrices	5
6	Write Matrix in CSV File	5
7	Main Function Implementation	6
8	Included Cuda Library	10
9	Defined Operations as Enums	10
10	Main Function Implementation	10
11	Matrix Operation Wrapper Function	10
12	Pointers to GPU Memory	11
13	Required GPU Memory Size	11
14	GPU Memory Allocation	11
15	Copying Data From CPU Memory to GPU Memory	11
16	Choosing Number of Threads and Block Size	11
17	Launching Appropriate Kernel Based on Passed Operation	12
18	Matrix Addition CUDA Kernel Function	12
19	Matrix Multiplication CUDA Kernel Function	13
20	Copying Data From GPU Memory to CPU Memory	13
21	Free Allocated GPU Memory	13

1 Introduction

This report contains code—and rigorous analysis of said code—written using normal C++ and using the Cuda framework in an effort to build a solid understanding on the differences between vector processors and ordinary processors. A program that calculates the matrix equation shown in 1 is run on both at different matrix sizes. Additionally, profiling has been used to observe each function’s timing to know what kind of operation consumes time in each version of the program. All code, profiling output files, and screenshots can be found in the project’s [Github repository](#).

$$R = C(A \times B + B \times A) \tag{1}$$

2 Prerequisites

This section contains the host system specifications used to run the two versions of the code, as well as a quick summary of the steps I followed in order to get my setup ready to compile Cuda code.

2.1 Host System Specifications

My setup is as shown in figure 1.

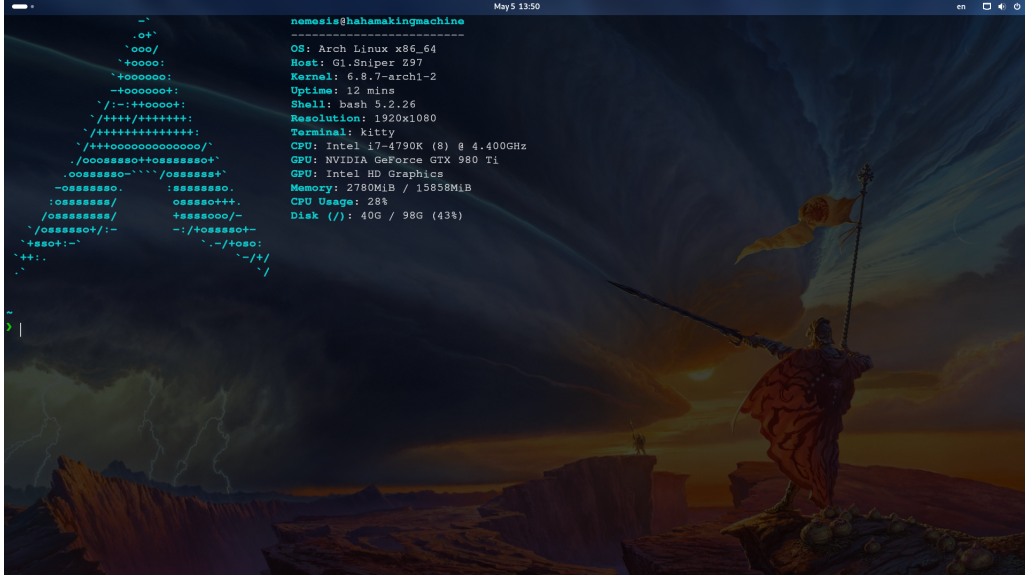


Figure 1: Host System Specifications

The GPU used is a GTX980ti: a flagship GPU released in 2015. The specifications relevant to this report are:

- Streaming Multiprocessor Count (SMC): 22 (each clocked at 1216MHz)
- CUDA Core Count (CCC): 2816 (SMC \times 128)
- Cache: 48 KB (L1 – per SM) & 3 MB (L2)
- Memory: 6GB of DDR5 (clocked at 1800 MHz)
- Memory Bus: 384 bit
- Memory Bandwidth: 345.6 GB/s

Understanding the GPU specifications is crucial for optimizing the block size and number of threads passed to the kernel—more on that later.

2.2 Installation Process

The process of installing Cuda was as follows:

1. Install cuda using

```
sudo pacman -S cuda
```

2. To make sure it was installed successfully:

```
nvcc --version
```

3. Export Cuda library directory into '.bashrc' as an environment variable.
4. Add the following library in the .cu program:

```
#include <cuda_runtime.h>
```

2.3 Compilation Settings

Since this lab's main purpose is to compare CPU and GPU performance in executing the same operation, CPU memory allocation has been avoided when defining the variables holding the matrix and the stack is instead used to make the comparison as fair as possible. Bare in mind that the default stack size would not be able to hold the matrices at large values of N which would result in a segmentation fault; so, the following command was used to extend the stack size—which is not a best practice but it is purely for testing purposes:

```
ulimit -s 4096000
```

When compiling the C++ version, I used gprof as a profiling tool to get a good estimate for how long does each function take to execute. So, the C++ code is compiled as follows:

```
g++ -pg nocuda.cpp -o nocuda
```

The program then have to at least be run once in order to extract timing information using the following command:

```
gprof nocuda gmon.out > analysis.txt
```

When compiling the Cuda version, nvcc is used as the compiler and, luckily, Nvidia's provided nvprof tool grants way more accurate timing information even without extra flags. To compile, simply run:

```
nvcc cuda.cu -o cuda
```

To analyse using nvprof, run:

```
time nvprof ./cuda
```

In the previous command, the shell command 'time' is used to give me the program's total runtime.

3 CPU Version of The Program

In this section, the CPU version of the program will be explained function-by-function. Additionally, runtime and profiling information will be observed at different values of N . ***Note that gprof is a bit inaccurate since its lowest time unit is 0.01s; ergo, at low values of N , information gathered by gprof will not be provided.***

3.1 Program Code

The program simply generates 3 square matrices—as well as 2 empty matrices: one as a temporary variable and one to store the final result—of size $N \times N$ and populates them with random numbers between 0 and 49. The generated matrices, as well as the matrix resulting from the operation are all written in comma separated (csv) files for debugging purposes.

```

1  #include <iostream>
2  #include <cstdlib>
3  #include <ctime>
4  #include <fstream>
5
6  using namespace std;
7
8  #define N 5000 // matrix row/col

```

Code Snippet 1: Included Libraries & Macro

```

1  void populateMatrix(unsigned int (&matrix) [N] [N])
2  {
3      for (int i = 0; i < N; i++)
4      {
5          for (int j = 0; j < N; j++)
6          {
7              matrix[i][j] = rand() % 50; // Generate random numbers
              // between 0 and 49 (inclusive)
8          }
9      }
10 }

```

Code Snippet 2: Function to Populate Matrices with Random Values

```

1  void printMatrix(const unsigned int (&matrix) [N] [N])
2  {
3      for (int i = 0; i < N; i++)
4      {
5          for (int j = 0; j < N; j++)
6          {
7              cout << matrix[i][j] << "\t";
8          }
9          cout << endl;
10     }
11 }

```

Code Snippet 3: Function to Print Matrix (For Testing Only)


```

1 void addMatrix(const unsigned int (&matrixA)[N][N], const unsigned int
2   (&matrixB)[N][N], unsigned int (&matrixC)[N][N])
3 {
4     for (int i = 0; i < N; i++)
5     {
6         for (int j = 0; j < N; j++)
7         {
8             matrixC[i][j] = matrixA[i][j] + matrixB[i][j];
9         }
10    }

```

Code Snippet 4: Function to Add Two Matrices

```

1 void multMatrix(const unsigned int (&matrixA)[N][N], const unsigned int
2   (&matrixB)[N][N], unsigned int (&matrixC)[N][N])
3 {
4     for (int i = 0; i < N; i++)
5     {
6         for (int j = 0; j < N; j++)
7         {
8             matrixC[i][j] = 0;
9             for (int k = 0; k < N; k++)
10            {
11                matrixC[i][j] += matrixA[i][k] * matrixB[k][j];
12            }
13        }
14    }

```

Code Snippet 5: unction to Multiply Two Matrices

```

1 void csvMatrix(const unsigned int (&matrix)[N][N], const char *filename)
2 {
3     std::ofstream file(filename);
4     for (int i = 0; i < N; i++)
5     {
6         for (int j = 0; j < N; j++)
7         {
8             file << matrix[i][j];
9             if (j < N - 1)
10            {
11                file << ",";
12            }
13        }
14        file << "\n";
15    }
16    file.close();

```

Code Snippet 6: Write Matrix in CSV File

```
1  int main()
2  {
3      // random number generation shenanigans
4      srand(time(NULL));
5
6      // define matrices
7      unsigned int matA[N][N];
8      unsigned int matB[N][N];
9      unsigned int matC[N][N];
10     unsigned int matRes[N][N];
11     unsigned int matTemp[N][N];
12
13     // populate matrix A && B
14     populateMatrix(matA);
15     populateMatrix(matB);
16     populateMatrix(matC);
17
18     // output matrix A && B as csv files for references
19     csvMatrix(matA, "MatrixA.csv");
20     csvMatrix(matB, "MatrixB.csv");
21     csvMatrix(matC, "MatrixC.csv");
22
23     // C * ((A * B) + (B * A))
24     multMatrix(matA, matB, matRes); // A * B = res
25
26     multMatrix(matB, matA, matTemp); // B * A = temp
27
28     addMatrix(matRes, matTemp, matTemp); // res + temp = temp
29
30     multMatrix(matC, matTemp, matRes); // C * temp = res
31
32     // output matrix C result for reference
33     csvMatrix(matRes, "Result.csv");
34
35     return 0;
36 }
```

Code Snippet 7: Main Function Implementation

3.2 Runtime & Profiling Results

In this section, runtime and profiling information will be observed for 6 values of N: 50, 100, 500, 1000, 2000, and 5000.

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    0m0.004s
user    0m0.003s
sys     0m0.001s
```

Figure 2: Runtime at N = 50

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    0m0.018s
user    0m0.018s
sys     0m0.000s
```

Figure 3: Runtime at N = 100

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    0m1.211s
user    0m1.203s
sys     0m0.007s
```

Figure 4: Runtime at N = 500

```
Each sample counts as 0.01 seconds.
%   cumulative   self           self         total
time  seconds    seconds   calls   ms/call  ms/call  name
97.95      1.10      1.10        3    365.70    365.70  multMatrix(unsig
 0.89      1.11      0.01        1     10.02     10.02  addMatrix(unsig
 0.00      1.11      0.00        4        0.00      0.00  csvMatrix(unsig
 0.00      1.11      0.00        3        0.00      0.00  populateMatrix(
```

Figure 5: Function Timing at N = 500 (gprof)

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    0m9.456s
user    0m9.423s
sys     0m0.010s
```

Figure 6: Runtime at N = 1000

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	s/call	s/call	name
98.45	9.08	9.08	3	3.03	3.03	multMatrix(unsig
0.11	9.09	0.01	4	0.00	0.00	csvMatrix(unsig
0.11	9.10	0.01	1	0.01	0.01	addMatrix(unsig
0.00	9.10	0.00	3	0.00	0.00	populateMatrix(

Figure 7: Function Timing at $N = 1000$ (gprof)

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    1m29.660s
user    1m29.257s
sys     0m0.175s
```

Figure 8: Runtime at $N = 2000$

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	s/call	s/call	name
99.30	87.95	87.95	3	29.32	29.32	multMatrix(unsig
0.03	87.98	0.03	4	0.01	0.01	csvMatrix(unsig
0.02	88.00	0.02	3	0.01	0.01	populateMatrix(
0.01	88.01	0.01	1	0.01	0.01	addMatrix(unsig

Figure 9: Function Timing at $N = 2000$ (gprof)

```
ComputerArchitecture_Playground/Lab_3 on main
> time ./nocuda

real    35m43.403s
user    35m39.282s
sys     0m0.613s
```

Figure 10: Runtime at $N = 5000$

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	s/call	s/call	name
99.46	1470.95	1470.95	3	490.32	490.32	multMatrix(unsig
0.01	1471.12	0.17	4	0.04	0.04	csvMatrix(unsig
0.01	1471.26	0.14	3	0.05	0.05	populateMatrix(u
0.00	1471.32	0.06	1	0.06	0.06	addMatrix(unsig
0.00	1471.35	0.03				_init

Figure 11: Function Timing at $N = 5000$ (gprof)

3.3 Observation

At low values of N , the CPU code is blazing fast and outperforms the GPU in everything. But, once N crosses the 500 mark, the performance shifts in favour of GPU because the `multMatrix` function simply takes too long to execute. Additionally, the time taken at values of N higher than 2000 is immense. Figure 12 shows the exponential growth that happens at very high values of N ¹.

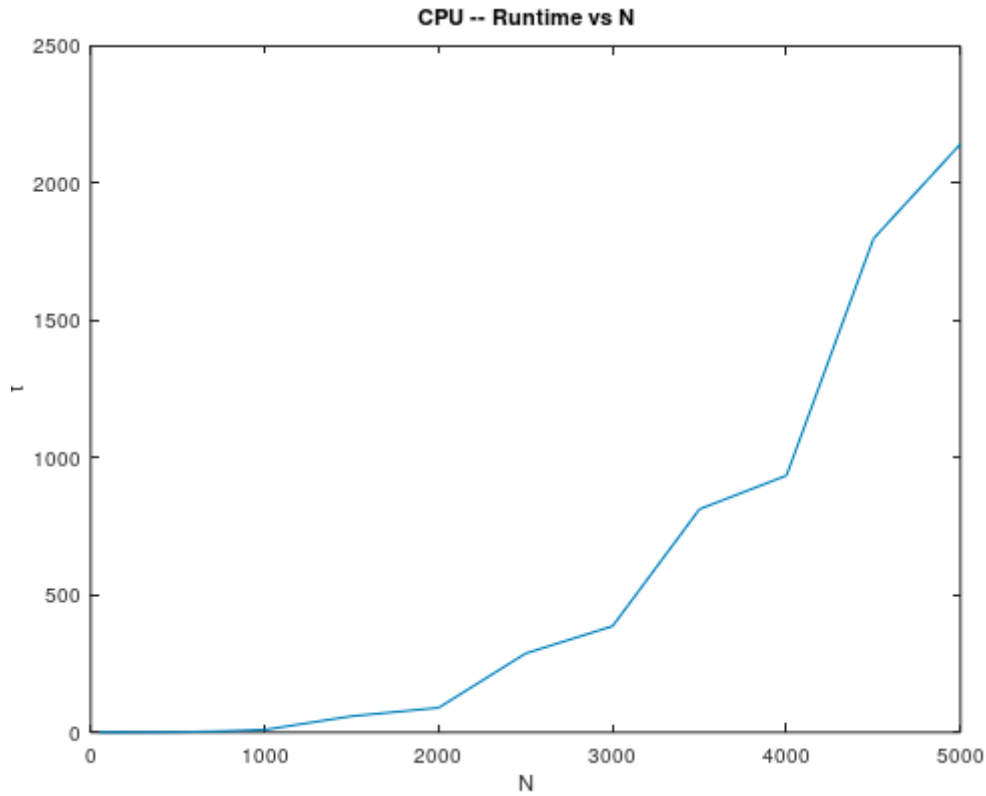


Figure 12: CPU – Runtime vs N

¹Graph was plotted using Matlab at $N = 50, 100, 250, 500, 750, 1000, 1500, 2000, 2500, 3000, 3500, 4000, 4500, 5000$.

4 GPU Version of The Program

In this section, the GPU version of the program will be explained function-by-function. Additionally, runtime and profiling information will be observed at different values of N.

4.1 Program Code

The program is mostly the same in many areas to its CPU counterpart—same `csvMatrix` and `populateMatrix`. But, the main difference is in how the GPU handles addition and multiplication, which will be explained later.

```
1 #include <cuda_runtime.h>
```

Code Snippet 8: Included Cuda Library

An addition to the code is determining the operation used based on 'enums' defined to ease code readability and simplify the used logic.

```
1 // available matrix operations
2 typedef enum
3 {
4     ADD = 1,
5     MUL = 2
6 } matOp;
```

Code Snippet 9: Defined Operations as Enums

The main change in the main function implementation is an entirely rewritten function to perform either addition or multiplication based on the passed enum operation. This function contains all the preparations for the GPU kernel functions for addition and multiplication.

```
1 // C * ((A * B) + (B * A))
2 matrixOperationCudaWrapper(matA, matB, matRes, MUL); // A * B = res
3
4 matrixOperationCudaWrapper(matB, matA, matTemp, MUL); // B * A = temp
5
6 matrixOperationCudaWrapper(matRes, matTemp, matTemp, ADD); // res +
   temp = temp
7
8 matrixOperationCudaWrapper(matC, matTemp, matRes, MUL); // C * temp
   = res
```

Code Snippet 10: Main Function Implementation

The CUDA wrapper function takes two 2D arrays, one 2D array to store the result, and an enum that tells the function which operation to execute.

```
1 void matrixOperationCudaWrapper(const unsigned int (&h_matrixA)[N][N],
   const unsigned int (&h_matrixB)[N][N], unsigned int
   (&h_matrixRes)[N][N], unsigned char operation)
```

Code Snippet 11: Matrix Operation Wrapper Function

The first thing this function does is to define pointers that will be used to store memory addresses.

```

1 // create pointers to gpu
2 unsigned int* d_cudaA = 0;
3 unsigned int* d_cudaB = 0;
4 unsigned int* d_cudaRes = 0;

```

Code Snippet 12: Pointers to GPU Memory

```

1 // defining size
2 size_t sizeInBytes = N * N * sizeof(unsigned int);

```

Code Snippet 13: Required GPU Memory Size

There are two ways to allocate GPU memory using CUDA: `cudaMalloc` and `cudaMallocManaged`. `cudaMalloc` was ultimately chosen because it was easier to understand its syntax and the benefits `cudaMallocManaged` brought was too advanced to understand.

```

1 // allocate memory in gpu
2 cudaMalloc((void**)(&d_cudaA), sizeInBytes);
3 cudaMalloc((void**)(&d_cudaB), sizeInBytes);
4 cudaMalloc((void**)(&d_cudaRes), sizeInBytes);

```

Code Snippet 14: GPU Memory Allocation

Here, the passed matrices to the function are copied to the allocated GPU memory. This copying process represents a huge overhead in the program and can be optimized by using `cudaMallocManaged` which eliminates the need to use `cudaMemcpy` completely.

```

1 // copy vectors into gpu
2 cudaMemcpy(d_cudaA, h_matrixA, sizeInBytes, cudaMemcpyHostToDevice);
3 cudaMemcpy(d_cudaB, h_matrixB, sizeInBytes, cudaMemcpyHostToDevice);

```

Code Snippet 15: Copying Data From CPU Memory to GPU Memory

The number of threads and block size are chosen based on the GPU's hardware. Through trial and error, the best performance was achieved at a thread count of 32 and a block size of 128. It is important to keep both numbers a power of two to optimize performance. The type `dim3` outputs The vector thread indices (X, Y, Z) suitable for the defined block size and thread count; this will be used in the kernel function.

```

1 // defining threads and block size
2 int threads = 32;
3 int blocks = 128;
4
5 // setting up kernel launch parameters
6 dim3 BLOCKS(blocks, blocks);
7 dim3 THREADS(threads, threads);

```

Code Snippet 16: Choosing Number of Threads and Block Size

Depending on the chosen operation, the corresponding kernel is launched by this obscurely complex syntax.

```

1 // launch kernel for chosen operation
2 if (operation == ADD)
3     matrixAddCUDA<<<BLOCKS, THREADS>>>(d_cudaA, d_cudaB, d_cudaRes);
4 else if (operation == MUL)
5     matrixMulCUDA<<<BLOCKS, THREADS>>>(d_cudaA, d_cudaB, d_cudaRes);
6 else
7 {
8     cout << "chotto matte!" << endl;
9     return; // do not continue this mess!
10 }

```

Code Snippet 17: Launching Appropriate Kernel Based on Passed Operation

The kernel function for addition is fairly simple. The vector thread indices will be used to index the memory locations to perform summation. The tricky part is to index a 2D array correctly using 1D array indexing. It is important to make sure that a thread operates within the array bounds.

```

1 __global__ void matrixAddCUDA(unsigned int* matrixA, unsigned int*
2   matrixB, unsigned int* matrixRes)
3 {
4     // Compute each thread's global row and column index
5     int rowIndex = blockIdx.y * blockDim.y + threadIdx.y;
6     int colIndex = blockIdx.x * blockDim.x + threadIdx.x;
7
8     if (rowIndex < N && colIndex < N)
9     {
10         // simply add
11         matrixRes[rowIndex * N + colIndex] = matrixA[rowIndex * N +
12         colIndex] + matrixB[rowIndex * N + colIndex];
13     }
14 }

```

Code Snippet 18: Matrix Addition CUDA Kernel Function

The kernel function for multiplication is a bit more complex. Like the previous function, it uses the vector thread indices to correctly index the arrays and perform a boundary check before executing any operation. It is important to initialize the result's element to 0 before accumulating the results of multiplying A and B.


```

1  __global__ void matrixMulCUDA(unsigned int* matrixA, unsigned int*
2  matrixB, unsigned int* matrixRes)
3  {
4      // Compute each thread's global row and column index
5      int rowIndex = blockIdx.y * blockDim.y + threadIdx.y;
6      int colIndex = blockIdx.x * blockDim.x + threadIdx.x;
7
8      // Iterate over row, and down column
9      if (rowIndex < N && colIndex < N)
10     {
11         matrixRes[rowIndex * N + colIndex] = 0;
12         for (int k = 0; k < N; k++)
13         {
14             // Accumulate results for a single element
15             matrixRes[rowIndex * N + colIndex] += matrixA[rowIndex * N +
16             k] * matrixB[k * N + colIndex];
17         }
18     }
19 }

```

Code Snippet 19: Matrix Multiplication CUDA Kernel Function

After the operation is done, the results are copied from GPU memory to CPU memory.

```

1  // copy result from gpu memory
2  cudaMemcpy(h_matrixRes, d_cudaRes, sizeInBytes,
3  cudaMemcpyDeviceToHost);

```

Code Snippet 20: Copying Data From GPU Memory to CPU Memory

Finally, free all the allocated GPU memory to avoid any memory leaks.

```

1  // free allocated gpu memory
2  cudaFree(d_cudaA);
3  cudaFree(d_cudaB);
4  cudaFree(d_cudaRes);

```

Code Snippet 21: Free Allocated GPU Memory

4.2 Runtime & Profiling Results

In this section, runtime and profiling information will be observed for 6 values of N at two thread counts (16 and 32): 50, 100, 500, 1000, 2000, and 5000.

```
ComputerArchitecture_Playground/Lab_3 on main [X!?:]
> time nvprof ./cuda
==4300== NVPROF is profiling process 4300, command: ./cuda
==4300== Profiling application: ./cuda
==4300== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 67.05% 43.170us    3 14.390us 13.888us 15.105us matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
17.89% 11.521us    8 1.4400us 1.4080us 1.5040us [CUDA memcpy HtoD]
11.48% 7.3920us    4 1.8480us 1.7600us 2.0800us [CUDA memcpy DtoH]
3.58% 2.3040us    1 2.3040us 2.3040us 2.3040us matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls: 96.66% 54.205ms 12 4.5171ms 1.5420us 53.973ms cudaMalloc
1.74% 976.11us    4 244.03us 6.6380us 953.73us cudaLaunchKernel
0.99% 555.81us 12 46.317us 4.9020us 409.69us cudaMemcpy
0.36% 201.28us 12 16.773us 1.7180us 65.051us cudaFree
0.20% 113.73us 114 997ns 73ns 47.850us cuDeviceGetAttribute
0.03% 14.160us 1 14.160us 14.160us 14.160us cuDeviceGetName
0.01% 6.6100us 1 6.6100us 6.6100us 6.6100us cuDeviceGetPCIBusId
0.00% 1.0710us 3 357ns 104ns 804ns cuDeviceGetCount
0.00% 457ns 2 228ns 110ns 347ns cuDeviceGet
0.00% 454ns 1 454ns 454ns 454ns cuDeviceTotalMem
0.00% 230ns 1 230ns 230ns 230ns cuModuleGetLoadingMode
0.00% 229ns 1 229ns 229ns 229ns cuDeviceGetUuid

real    0m0.400s
user    0m0.022s
sys     0m0.102s
```

Figure 13: Program Timing Information at N = 50 (Thread Count of 16)

```
ComputerArchitecture_Playground/Lab_3 on main [X!?:]
> time nvprof ./cuda
==4480== NVPROF is profiling process 4480, command: ./cuda
==4480== Profiling application: ./cuda
==4480== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 58.92% 86.370us    3 28.790us 28.160us 29.697us matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
27.90% 40.896us    8 5.1120us 5.0880us 5.1520us [CUDA memcpy HtoD]
11.48% 16.832us    4 4.2080us 4.0000us 4.4160us [CUDA memcpy DtoH]
1.70% 2.4960us    1 2.4960us 2.4960us 2.4960us matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls: 95.34% 44.635ms 12 3.7196ms 1.5880us 44.404ms cudaMalloc
3.18% 1.4873ms 114 13.046us 117ns 1.3712ms cuDeviceGetAttribute
0.73% 341.18us 12 28.431us 10.714us 99.393us cudaMemcpy
0.42% 198.05us 12 16.504us 1.8380us 62.790us cudaFree
0.25% 119.03us 4 29.757us 6.5100us 96.605us cudaLaunchKernel
0.05% 22.120us 1 22.120us 22.120us 22.120us cuDeviceGetName
0.02% 9.4010us 1 9.4010us 9.4010us 9.4010us cuDeviceGetPCIBusId
0.00% 1.3190us 3 439ns 184ns 950ns cuDeviceGetCount
0.00% 706ns 2 353ns 159ns 547ns cuDeviceGet
0.00% 507ns 1 507ns 507ns 507ns cuDeviceTotalMem
0.00% 388ns 1 388ns 388ns 388ns cuModuleGetLoadingMode
0.00% 316ns 1 316ns 316ns 316ns cuDeviceGetUuid

real    0m0.221s
user    0m0.025s
sys     0m0.076s
```

Figure 14: Program Timing Information at N = 100 (Thread Count of 16)

```
ComputerArchitecture_Playground/Lab_3 on main [X!?:]
> time nvprof ./cuda
==4665== NVPROF is profiling process 4665, command: ./cuda
==4665== Profiling application: ./cuda
==4665== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 88.82% 7.7082ms    3 2.5694ms 2.5515ms 2.5835ms matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
7.42% 644.11us    8 80.514us 79.746us 82.306us [CUDA memcpy HtoD]
3.66% 317.32us    4 79.329us 79.138us 79.650us [CUDA memcpy DtoH]
0.11% 9.1840us    1 9.1840us 9.1840us 9.1840us matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls: 78.60% 43.353ms 12 3.6127ms 43.213us 42.774ms cudaMalloc
19.16% 10.569ms 12 880.75us 102.58us 3.0784ms cudaMemcpy
1.70% 935.99us 12 77.998us 44.118us 89.788us cudaFree
0.30% 163.62us 4 40.904us 8.4900us 124.50us cudaLaunchKernel
0.20% 112.63us 114 987ns 71ns 47.385us cuDeviceGetAttribute
0.02% 13.199us 1 13.199us 13.199us 13.199us cuDeviceGetName
0.01% 7.0780us 1 7.0780us 7.0780us 7.0780us cuDeviceGetPCIBusId
0.00% 1.3600us 3 453ns 110ns 1.1350us cuDeviceGetCount
0.00% 531ns 2 265ns 140ns 391ns cuDeviceGet
0.00% 444ns 1 444ns 444ns 444ns cuDeviceTotalMem
0.00% 231ns 1 231ns 231ns 231ns cuDeviceGetUuid
0.00% 218ns 1 218ns 218ns 218ns cuModuleGetLoadingMode

real    0m0.365s
user    0m0.088s
sys     0m0.074s
```

Figure 15: Program Timing Information at N = 500 (Thread Count of 16)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==4841== NVPROF is profiling process 4841, command: ./cuda
==4841== Profiling application: ./cuda
==4841== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 87.77% 53.223ms      3      17.741ms  17.603ms  17.998ms  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
                7.60% 4.6054ms      8      575.79us  539.09us  682.45us  [CUDA memcpy HtoD]
                4.55% 2.7614ms      4      690.35us  596.94us  797.46us  [CUDA memcpy DtoH]
                0.08% 48.065us      1      48.065us  48.065us  48.065us  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:      54.82% 62.542ms      12      5.2119ms  599.07us  18.908ms  cudaMemcpy
                42.50% 48.491ms      12      4.0409ms  43.158us  47.327ms  cudaMalloc
                2.32% 2.6482ms      12      220.68us  80.914us  281.63us  cudaFree
                0.20% 228.79us      4      57.197us  26.745us  138.31us  cudaLaunchKernel
                0.13% 151.14us      114      1.3250us  111ns    65.033us  cuDeviceGetAttribute
                0.01% 14.995us      1      14.995us  14.995us  14.995us  cuDeviceGetName
                0.01% 7.3960us      1      7.3960us  7.3960us  7.3960us  cuDeviceGetPCIBusId
                0.00% 1.2520us      3      417ns    152ns    867ns    cuDeviceGetCount
                0.00% 777ns      2      388ns    138ns    639ns    cuDeviceGet
                0.00% 576ns      1      576ns    576ns    576ns    cuDeviceTotalMem
                0.00% 383ns      1      383ns    383ns    383ns    cuModuleGetLoadingMode
                0.00% 280ns      1      280ns    280ns    280ns    cuDeviceGetUuid

real    0m0.601s
user    0m0.309s
sys     0m0.087s

```

Figure 16: Program Timing Information at N = 1000 (Thread Count of 16)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==5006== NVPROF is profiling process 5006, command: ./cuda
==5006== Profiling application: ./cuda
==5006== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 86.68% 218.70ms      3      72.900ms  70.105ms  78.336ms  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
                8.07% 20.361ms      8      2.5451ms  2.4427ms  2.7509ms  [CUDA memcpy HtoD]
                5.18% 13.066ms      4      3.2664ms  2.5716ms  4.0852ms  [CUDA memcpy DtoH]
                0.07% 185.51us      1      185.51us  185.51us  185.51us  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:      83.99% 254.01ms      12      21.168ms  2.5143ms  82.581ms  cudaMemcpy
                13.38% 40.452ms      12      3.3710ms  45.970us  37.421ms  cudaMalloc
                2.51% 7.5998ms      12      633.32us  114.24us  900.01us  cudaFree
                0.08% 255.89us      4      63.972us  31.279us  149.41us  cudaLaunchKernel
                0.03% 98.207us      114      861ns    72ns    41.795us  cuDeviceGetAttribute
                0.00% 12.147us      1      12.147us  12.147us  12.147us  cuDeviceGetName
                0.00% 6.8830us      1      6.8830us  6.8830us  6.8830us  cuDeviceGetPCIBusId
                0.00% 792ns      3      264ns    104ns    559ns    cuDeviceGetCount
                0.00% 459ns      2      229ns    74ns    385ns    cuDeviceGet
                0.00% 380ns      1      380ns    380ns    380ns    cuDeviceTotalMem
                0.00% 245ns      1      245ns    245ns    245ns    cuDeviceGetUuid
                0.00% 233ns      1      233ns    233ns    233ns    cuModuleGetLoadingMode

real    0m1.559s
user    0m1.146s
sys     0m0.110s

```

Figure 17: Program Timing Information at N = 2000 (Thread Count of 16)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==5236== NVPROF is profiling process 5236, command: ./cuda
==5236== Profiling application: ./cuda
==5236== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 96.67% 5.93430s      3      1.97810s  1.97639s  1.98026s  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
                2.03% 124.76ms      8      15.595ms  15.224ms  16.522ms  [CUDA memcpy HtoD]
                1.28% 78.686ms      4      19.672ms  15.681ms  24.180ms  [CUDA memcpy DtoH]
                0.02% 1.1758ms      1      1.1758ms  1.1758ms  1.1758ms  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:      98.37% 6.14064s      12      511.72ms  15.333ms  2.00212s  cudaMemcpy
                1.07% 66.864ms      12      5.5720ms  64.555us  44.219ms  cudaMalloc
                0.55% 34.605ms      12      2.8837ms  137.46us  4.2751ms  cudaFree
                0.00% 286.64us      4      71.659us  27.984us  194.10us  cudaLaunchKernel
                0.00% 115.26us      114      1.0110us  75ns    49.556us  cuDeviceGetAttribute
                0.00% 11.585us      1      11.585us  11.585us  11.585us  cuDeviceGetName
                0.00% 6.6900us      1      6.6900us  6.6900us  6.6900us  cuDeviceGetPCIBusId
                0.00% 843ns      3      281ns    146ns    532ns    cuDeviceGetCount
                0.00% 531ns      1      531ns    531ns    531ns    cuDeviceTotalMem
                0.00% 456ns      2      228ns    90ns    366ns    cuDeviceGet
                0.00% 228ns      1      228ns    228ns    228ns    cuDeviceGetUuid
                0.00% 225ns      1      225ns    225ns    225ns    cuModuleGetLoadingMode

real    0m12.674s
user    0m11.270s
sys     0m0.445s

```

Figure 18: Program Timing Information at N = 5000 (Thread Count of 16)

```
ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==18316== NVPROF is profiling process 18316, command: ./cuda
==18316== Profiling application: ./cuda
==18316== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	73.87%	876.12us	3	292.04us	291.37us	292.87us	matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
	24.52%	290.76us	1	290.76us	290.76us	290.76us	matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
	0.97%	11.520us	8	1.4400us	1.4080us	1.4720us	[CUDA memcpy HtoD]
	0.64%	7.5530us	4	1.8880us	1.7600us	2.2400us	[CUDA memcpy DtoH]
API calls:	96.24%	44.556ms	12	3.7130ms	1.7090us	44.323ms	cudaMalloc
	2.80%	1.2941ms	12	107.84us	7.0840us	308.03us	cudaMemcpy
	0.43%	201.24us	12	16.769us	1.6970us	63.574us	cudaFree
	0.25%	114.61us	4	28.651us	6.7540us	91.930us	cudaLaunchKernel
	0.24%	109.68us	114	962ns	75ns	46.394us	cuDeviceGetAttribute
	0.02%	11.510us	1	11.510us	11.510us	11.510us	cuDeviceGetName
	0.01%	6.8040us	1	6.8040us	6.8040us	6.8040us	cuDeviceGetPCIBusId
	0.00%	935ns	3	311ns	116ns	651ns	cuDeviceGetCount
	0.00%	519ns	2	259ns	134ns	385ns	cuDeviceGet
	0.00%	473ns	1	473ns	473ns	473ns	cuDeviceTotalMem
	0.00%	296ns	1	296ns	296ns	296ns	cuModuleGetLoadingMode
	0.00%	142ns	1	142ns	142ns	142ns	cuDeviceGetUuid

```

real    0m0.230s
user    0m0.023s
sys     0m0.088s

```

Figure 19: Program Timing Information at N = 50 (Thread Count of 32)

```
ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==18501== NVPROF is profiling process 18501, command: ./cuda
==18501== Profiling application: ./cuda
==18501== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	72.24%	906.84us	3	302.28us	301.38us	302.79us	matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
	23.17%	290.89us	1	290.89us	290.89us	290.89us	matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
	3.26%	40.962us	8	5.1200us	5.0890us	5.1840us	[CUDA memcpy HtoD]
	1.32%	16.578us	4	4.1440us	4.0320us	4.3530us	[CUDA memcpy DtoH]
API calls:	96.10%	45.939ms	12	3.8282ms	1.7320us	45.707ms	cudaMalloc
	2.94%	1.4074ms	12	117.28us	10.368us	332.62us	cudaMemcpy
	0.44%	209.19us	12	17.432us	1.8500us	64.018us	cudaFree
	0.25%	119.05us	4	29.763us	7.0810us	95.293us	cudaLaunchKernel
	0.23%	110.01us	114	964ns	69ns	47.447us	cuDeviceGetAttribute
	0.02%	8.8700us	1	8.8700us	8.8700us	8.8700us	cuDeviceGetName
	0.01%	7.1220us	1	7.1220us	7.1220us	7.1220us	cuDeviceGetPCIBusId
	0.00%	865ns	3	288ns	99ns	645ns	cuDeviceGetCount
	0.00%	545ns	2	272ns	87ns	458ns	cuDeviceGet
	0.00%	489ns	1	489ns	489ns	489ns	cuDeviceTotalMem
	0.00%	238ns	1	238ns	238ns	238ns	cuModuleGetLoadingMode
	0.00%	215ns	1	215ns	215ns	215ns	cuDeviceGetUuid

```

real    0m0.221s
user    0m0.022s
sys     0m0.077s

```

Figure 20: Program Timing Information at N = 100 (Thread Count of 32)

```
ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==19080== NVPROF is profiling process 19080, command: ./cuda
==19080== Profiling application: ./cuda
==19080== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	82.04%	5.7252ms	3	1.9084ms	1.8435ms	2.0255ms	matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
	9.20%	641.91us	8	80.238us	80.066us	80.514us	[CUDA memcpy HtoD]
	4.54%	316.71us	4	79.177us	79.106us	79.265us	[CUDA memcpy DtoH]
	4.23%	294.92us	1	294.92us	294.92us	294.92us	matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:	82.09%	45.365ms	12	3.7804ms	42.782us	44.789ms	cudaMalloc
	15.73%	8.6945ms	12	724.54us	71.539us	2.3589ms	cudaMemcpy
	1.68%	926.20us	12	77.183us	44.960us	93.592us	cudaFree
	0.27%	147.93us	4	36.981us	8.0840us	121.23us	cudaLaunchKernel
	0.20%	110.55us	114	969ns	72ns	47.870us	cuDeviceGetAttribute
	0.02%	9.8620us	1	9.8620us	9.8620us	9.8620us	cuDeviceGetName
	0.01%	7.2150us	1	7.2150us	7.2150us	7.2150us	cuDeviceGetPCIBusId
	0.00%	762ns	3	254ns	102ns	525ns	cuDeviceGetCount
	0.00%	485ns	2	242ns	107ns	378ns	cuDeviceGet
	0.00%	449ns	1	449ns	449ns	449ns	cuDeviceTotalMem
	0.00%	244ns	1	244ns	244ns	244ns	cuModuleGetLoadingMode
	0.00%	221ns	1	221ns	221ns	221ns	cuDeviceGetUuid

```

real    0m0.374s
user    0m0.088s
sys     0m0.086s

```

Figure 21: Program Timing Information at N = 500 (Thread Count of 32)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==19505== NVPROF is profiling process 19505, command: ./cuda
==19505== Profiling application: ./cuda
==19505== Profiling result:
   Type  Time(%)   Time     Calls   Avg       Min       Max  Name
GPU activities: 81.91%  33.988ms    3  11.329ms  11.023ms  11.744ms  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
               10.83%  4.4933ms    8   561.67us  526.70us  658.29us  [CUDA memcpy HtoD]
               6.51%  2.6998ms    4   674.95us  570.35us  810.71us  [CUDA memcpy DtoH]
               0.76%   313.96us    1   313.96us  313.96us  313.96us  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:    50.23%  46.553ms   12   3.8794ms  44.931us  45.382ms  cudaMalloc
               46.57%  43.159ms   12   3.5966ms  581.54us  12.812ms  cudaMemcpy
               2.78%   2.5737ms   12   214.47us  59.938us  282.19us  cudaFree
               0.21%   191.81us    4   47.953us  12.391us  125.52us  cudaLaunchKernel
               0.18%   163.63us   114   1.4350us    117ns   64.025us  cuDeviceGetAttribute
               0.02%   23.123us    1   23.123us  23.123us  23.123us  cuDeviceGetName
               0.01%   10.452us    1   10.452us  10.452us  10.452us  cuDeviceGetPCIBusId
               0.00%   1.6960us    3    565ns    173ns   1.3330us  cuDeviceGetCount
               0.00%    698ns    2    349ns    164ns    534ns  cuDeviceGet
               0.00%    569ns    1    569ns    569ns    569ns  cuDeviceTotalMem
               0.00%    374ns    1    374ns    374ns    374ns  cuModuleGetLoadingMode
               0.00%    311ns    1    311ns    311ns    311ns  cuDeviceGetUuid

real    0m0.584s
user    0m0.275s
sys     0m0.089s

```

Figure 22: Program Timing Information at N = 1000 (Thread Count of 32)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==19853== NVPROF is profiling process 19853, command: ./cuda
==19853== Profiling application: ./cuda
==19853== Profiling result:
   Type  Time(%)   Time     Calls   Avg       Min       Max  Name
GPU activities: 87.14%  218.54ms    3  72.846ms  70.569ms  75.371ms  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
               7.81%   19.577ms    8   2.4471ms  2.3605ms  2.7227ms  [CUDA memcpy HtoD]
               4.91%   12.313ms    4   3.0783ms  2.4502ms  3.8287ms  [CUDA memcpy DtoH]
               0.14%    347.72us    1   347.72us  347.72us  347.72us  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:    83.22%  252.46ms   12   21.038ms  2.4909ms  79.256ms  cudaMemcpy
               14.11%  42.812ms   12   3.5677ms  46.080us  39.727ms  cudaMalloc
               2.54%   7.6922ms   12   641.02us  104.44us  957.35us  cudaFree
               0.09%   280.05us    4   70.013us  29.073us  186.00us  cudaLaunchKernel
               0.03%    96.141us   114    843ns    77ns   41.132us  cuDeviceGetAttribute
               0.00%   10.740us    1   10.740us  10.740us  10.740us  cuDeviceGetName
               0.00%   7.2890us    1   7.2890us  7.2890us  7.2890us  cuDeviceGetPCIBusId
               0.00%    886ns    3    295ns    94ns   645ns  cuDeviceGetCount
               0.00%    503ns    1    503ns    503ns    503ns  cuDeviceTotalMem
               0.00%    250ns    2    250ns    100ns   401ns  cuDeviceGet
               0.00%    296ns    1    296ns    296ns    296ns  cuModuleGetLoadingMode
               0.00%    233ns    1    233ns    233ns    233ns  cuDeviceGetUuid

real    0m1.518s
user    0m1.133s
sys     0m0.090s

```

Figure 23: Program Timing Information at N = 2000 (Thread Count of 32)

```

ComputerArchitecture_Playground/Lab_3 on main [X!?]
> time nvprof ./cuda
==20774== NVPROF is profiling process 20774, command: ./cuda
==20774== Profiling application: ./cuda
==20774== Profiling result:
   Type  Time(%)   Time     Calls   Avg       Min       Max  Name
GPU activities: 91.92%  2.34635s    3  782.12ms  770.43ms  793.02ms  matrixMulCUDA(unsigned int*, unsigned int*, unsigned int*)
               4.87%   124.38ms    8   15.548ms  15.259ms  16.584ms  [CUDA memcpy HtoD]
               3.18%   81.145ms    4   20.286ms  15.640ms  25.112ms  [CUDA memcpy DtoH]
               0.03%   791.67us    1   791.67us  791.67us  791.67us  matrixAddCUDA(unsigned int*, unsigned int*, unsigned int*)
API calls:    96.15%  2.55463s   12   212.89ms  15.326ms  818.46ms  cudaMemcpy
               2.53%   67.259ms   12   5.6050ms  67.490us  44.622ms  cudaMalloc
               1.30%   34.642ms   12   2.8868ms  141.28us  4.3188ms  cudaFree
               0.01%   291.16us    4   72.788us  29.494us  194.73us  cudaLaunchKernel
               0.00%   109.03us   114    956ns    74ns   46.894us  cuDeviceGetAttribute
               0.00%   9.8600us    1   9.8600us  9.8600us  9.8600us  cuDeviceGetName
               0.00%   6.7940us    1   6.7940us  6.7940us  6.7940us  cuDeviceGetPCIBusId
               0.00%    990ns    3    330ns    101ns    700ns  cuDeviceGetCount
               0.00%    499ns    2    249ns    106ns    393ns  cuDeviceGet
               0.00%    465ns    1    465ns    465ns    465ns  cuDeviceTotalMem
               0.00%    236ns    1    236ns    236ns    236ns  cuDeviceGetUuid
               0.00%    224ns    1    224ns    224ns    224ns  cuModuleGetLoadingMode

real    0m8.916s
user    0m7.682s
sys     0m0.394s

```

Figure 24: Program Timing Information at N = 5000 (Thread Count of 32)

4.3 Observation

At low values of N , the GPU code is ridiculously slow: the amount of time taken to copy data between GPU and CPU memory outweighs any gains in execution time of actual operations. Once N passes the 500 mark, actual gains start to appear. At values of N higher than 2000, the performance difference is staggering and the memory overhead is outshined by the speedup obtained in the multiplication operation. Figures 25 and 26 show the runtime at different values of N^2 at thread count 16 and 32 respectively. It is noticed that the slope of the second graph is more forgiving and that the runtime is considerably lower at the higher thread count.

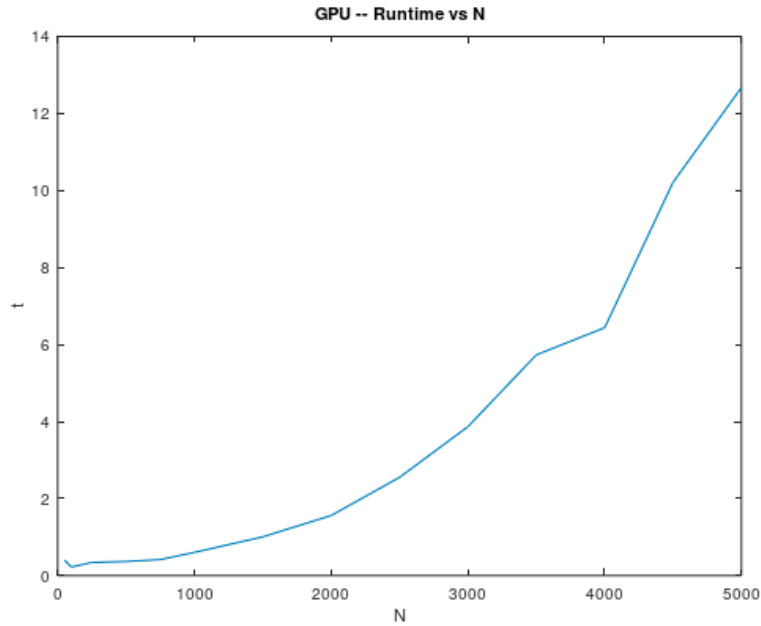


Figure 25: GPU – Runtime vs N (Thread Count of 16)

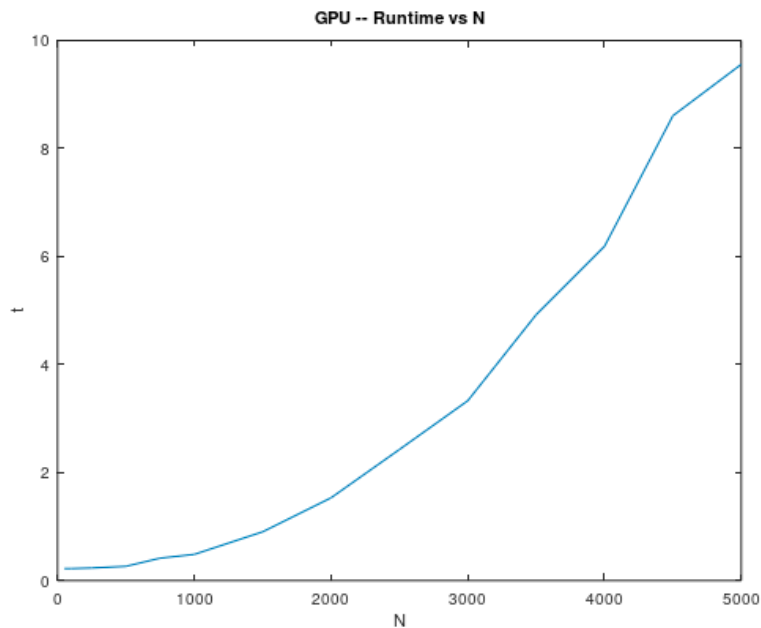


Figure 26: GPU – Runtime vs N (Thread Count of 32)

²Graph One and graph two were plotted using Matlab at $N = 50, 100, 250, 500, 750, 1000, 1500, 2000, 2500, 3000, 3500, 4000, 4500, 5000$.